

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: data = pd.read_csv('E:/New folder/sonar.all-data.csv')
data.head()
```

```
Out[2]:
```

	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	0.2111	...	0.0027	0.0065	0.0159
0	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	0.2872	...	0.0084	0.0089	0.0048
1	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	0.6194	...	0.0232	0.0166	0.0095
2	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	0.1264	...	0.0121	0.0036	0.0150
3	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	0.4459	...	0.0031	0.0054	0.0105
4	0.0286	0.0453	0.0277	0.0174	0.0384	0.0990	0.1201	0.1833	0.2105	0.3039	...	0.0045	0.0014	0.0038

5 rows × 61 columns

```
In [3]: data.isna().sum()
```

```
Out[3]: 0.0200    0
0.0371    0
0.0428    0
0.0207    0
0.0954    0
..
0.0180    0
0.0084    0
0.0090    0
0.0032    0
R         0
Length: 61, dtype: int64
```

There are 60 features, from which we can predict whether the object is rock or mine

```
In [4]: data['R'].value_counts()
```

```
Out[4]: M    111
R      96
Name: R, dtype: int64
```

```
In [5]: data_R = data[data['R']=='R']
data_M = data[data['R']=='M']
```

```
In [6]: data_R.head(5)
```

```
Out[6]:
```

	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	0.2111	...	0.0027	0.0065	0.0159
0	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	0.2872	...	0.0084	0.0089	0.0048
1	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	0.6194	...	0.0232	0.0166	0.0095
2	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	0.1264	...	0.0121	0.0036	0.0150
3	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	0.4459	...	0.0031	0.0054	0.0105
4	0.0286	0.0453	0.0277	0.0174	0.0384	0.0990	0.1201	0.1833	0.2105	0.3039	...	0.0045	0.0014	0.0038

5 rows × 61 columns

```
In [7]: data_M.head(5)
```


```
Out[7]:
```

	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	0.2111	...	0.0027	0.0065	0.00
96	0.0491	0.0279	0.0592	0.1270	0.1772	0.1908	0.2217	0.0768	0.1246	0.2028	...	0.0081	0.0129	0.00
97	0.1313	0.2339	0.3059	0.4264	0.4010	0.1791	0.1853	0.0055	0.1929	0.2231	...	0.0362	0.0210	0.00
98	0.0201	0.0423	0.0554	0.0783	0.0620	0.0871	0.1201	0.2707	0.1206	0.0279	...	0.0191	0.0182	0.00
99	0.0629	0.1065	0.1526	0.1229	0.1437	0.1190	0.0884	0.0907	0.2107	0.3597	...	0.0089	0.0262	0.00
100	0.0335	0.0134	0.0696	0.1180	0.0348	0.1180	0.1948	0.1607	0.3036	0.4372	...	0.0244	0.0232	0.00

5 rows × 61 columns


```
In [8]: plt.scatter(data['0.0200'],data['0.0371'])
```

```
Out[8]: <matplotlib.collections.PathCollection at 0x2339a312fc8>
```




```
In [9]: plt.scatter(data['0.0428'],data['0.0207'])
```

```
Out[9]: <matplotlib.collections.PathCollection at 0x2339a312b48>
```



```
In [10]: plt.plot(data['0.0200'], label='0.0200')
plt.plot(data['0.0428'], label='0.0428')
plt.plot(data['0.0207'], label='0.0207')
plt.legend()
```

Out[10]: <matplotlib.legend.Legend at 0x2339a44b908>



```
In [11]: from sklearn.model_selection import train_test_split
X=data.drop('R', axis=1)
y=data['R']
```

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(X,y, train_size=0.8, random_
state=29)
```

```
In [13]: print(X_train.shape)
print(X_test.shape)
```

(165, 60)
(42, 60)

Classification

KNN

```
In [14]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [15]: knn = KNeighborsClassifier()
```

```
In [16]: knn.fit(X=X_train, y=y_train)
```

Out[16]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
metric_params=None, n_jobs=None, n_neighbors=5, p=2,
weights='uniform')

```
In [17]: knn.score(X_test, y_test)
```


Out[17]: 0.7619047619047619

Accuracy: 76.2%

```
In [18]: neighbors = np.arange(1,9)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))

for i,k in enumerate(neighbors):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    train_accuracy[i] = knn.score(X_train,y_train)
    test_accuracy[i] = knn.score(X_test, y_test)

plt.plot(neighbors, test_accuracy, label='Testing Accuracy')
plt.plot(neighbors, train_accuracy, label='Training Accuracy')
plt.legend()
plt.show()
```



Maximum accuracy is obtained when k=3 or 4

```
In [19]: knn = KNeighborsClassifier(n_neighbors=4)
knn.fit(X=X_train, y=y_train)
knn.score(X_test, y_test)
```

Out[19]: 0.7857142857142857

Accuracy: 78.57%

```
In [20]: from sklearn.metrics import confusion_matrix
y_pred = knn.predict(X_test)
confusion_matrix(y_test, y_pred)
```

Out[20]: array([[23, 5],
 [4, 10]], dtype=int64)

9 wrong predictions out of 42

```
In [21]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
M	0.85	0.82	0.84	28
R	0.67	0.71	0.69	14
accuracy			0.79	42
macro avg	0.76	0.77	0.76	42
weighted avg	0.79	0.79	0.79	42

Decision Tree

```
In [22]: from sklearn.tree import DecisionTreeClassifier
         clf = DecisionTreeClassifier()
         clf.fit(X_train, y_train)
         clf.score(X_test, y_test)
```

Out[22]: 0.9047619047619048

Accuracy: 90.47%

```
In [23]: from sklearn.metrics import confusion_matrix
         y_pred = clf.predict(X_test)
         confusion_matrix(y_test, y_pred)
```

Out[23]: array([[27, 1],
 [3, 11]], dtype=int64)

4 out of 42 were predicted wrong

```
In [24]: from sklearn.metrics import classification_report
         print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
M	0.90	0.96	0.93	28
R	0.92	0.79	0.85	14
accuracy			0.90	42
macro avg	0.91	0.88	0.89	42
weighted avg	0.91	0.90	0.90	42

SVM

```
In [25]: from sklearn.svm import SVC
         svm = SVC()
         svm.fit(X_train, y_train)
         svm.score(X_test, y_test)
```

C:\Users\SR1407SM1106\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
 "avoid this warning.", FutureWarning)

Out[25]: 0.6428571428571429

```
In [26]: from sklearn.metrics import confusion_matrix
         y_pred = svm.predict(X_test)
         confusion_matrix(y_test, y_pred)
```

Out[26]: array([[21, 7],
 [8, 6]], dtype=int64)

```
In [27]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
M	0.72	0.75	0.74	28
R	0.46	0.43	0.44	14
accuracy			0.64	42
macro avg	0.59	0.59	0.59	42
weighted avg	0.64	0.64	0.64	42

Accuracy: 64.28%

15 were predicted wrong out of 42

Standardising data

```
In [42]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_p = scaler.fit_transform(X)
X_p
```

```
Out[42]: array([[ 0.70018948,  0.42042142,  1.0529498 , ..., -0.4709383 ,
                 -0.44268846, -0.42246083],
                [-0.13089402,  0.59942737,  1.71912994, ...,  1.30656071,
                 0.25299833,  0.25405324],
                [-0.83579208, -0.64754631,  0.48045125, ..., -0.54822087,
                 -0.63683361,  1.03005467],
                ...,
                [ 1.00042384,  0.15949749, -0.67235266, ...,  0.90469137,
                 -0.0382194 , -0.68112798],
                [ 0.0475061 , -0.09535845,  0.13434985, ..., -0.00724291,
                 -0.70154866, -0.34287094],
                [-0.13959647, -0.06501846, -0.78685237, ..., -0.67187297,
                 -0.2970796 ,  0.99025973]])
```

PCA

```
In [43]: from sklearn.decomposition import PCA
pca = PCA()
X_pca = pca.fit_transform(X_p)
```

```
In [44]: print(pca.explained_variance_ratio_)
```

```
[2.03897950e-01 1.89405455e-01 8.56018085e-02 5.68862840e-02
 5.01317021e-02 4.05567677e-02 3.27272864e-02 3.03529373e-02
 2.57210787e-02 2.49016628e-02 2.07971452e-02 1.90546557e-02
 1.75405558e-02 1.54132664e-02 1.43177628e-02 1.35130771e-02
 1.23002112e-02 1.11555925e-02 1.03282502e-02 9.79584100e-03
 9.31755378e-03 8.82245274e-03 8.41377607e-03 7.75326916e-03
 7.24358160e-03 6.98525179e-03 6.05789229e-03 5.32915402e-03
 5.25475666e-03 4.92103776e-03 4.77289354e-03 4.36932921e-03
 3.68829839e-03 3.26860405e-03 3.09930505e-03 3.03922888e-03
 2.87559886e-03 2.44611335e-03 2.19463709e-03 2.06394408e-03
 1.85188563e-03 1.58641602e-03 1.34933280e-03 1.24743835e-03
 1.01715873e-03 9.48405524e-04 9.01965543e-04 7.07573125e-04
 5.59258464e-04 5.23852009e-04 4.82666208e-04 4.46857015e-04
 3.80233202e-04 3.55950060e-04 3.20623754e-04 2.68355115e-04
 2.48356050e-04 1.92144557e-04 1.87217896e-04 1.08340357e-04]
```

```
In [45]: print(pca.singular_values_)
```

```
[50.32308163 48.50170877 32.60635616 26.5805878 24.95266999 22.44359719
 20.16117301 19.41606247 17.87332641 17.58632002 16.0717312 15.38371942
 14.759868 13.83592313 13.33516455 12.95501516 12.35996051 11.77083084
 11.32593784 11.03015617 10.75750984 10.46780125 10.22248007 9.8130323
 9.48500308 9.3143345 8.67404301 8.13560648 8.07861856 7.81788264
 7.69930762 7.36661855 6.76820996 6.37150393 6.20430244 6.14387685
 5.97619761 5.51187153 5.22086129 5.06302138 4.79587526 4.43883847
 4.09374076 3.93613825 3.55430885 3.43208342 3.34700046 2.96446592
 2.63552464 2.55073361 2.44841057 2.35583618 2.17313055 2.10259357
 1.99553176 1.8256425 1.75629785 1.54480918 1.52487582 1.1599945 ]
```

KNN

```
In [46]: from sklearn.model_selection import train_test_split
X_train_pca, X_test_pca, y_train_pca, y_test_pca = train_test_split(X_pca, y, te
st_size=0.2, random_state=1)
X_train_pca.shape
```

```
Out[46]: (165, 60)
```

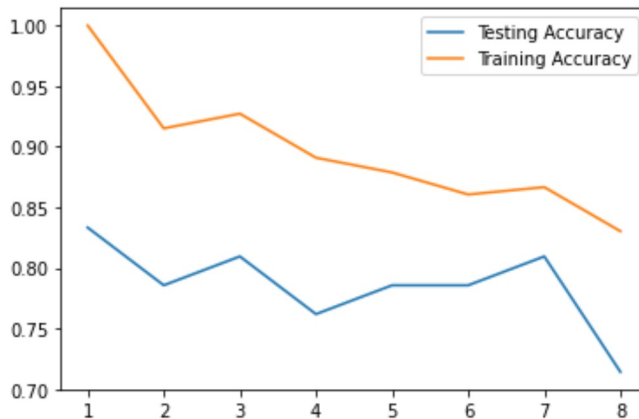
```
In [47]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier()
knn.fit(X=X_train_pca, y=y_train_pca)
knn.score(X_test_pca, y_test_pca)
```

```
Out[47]: 0.7857142857142857
```

```
In [48]: neighbors = np.arange(1,9)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))

for i,k in enumerate(neighbors):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train_pca, y_train_pca)
    train_accuracy[i] = knn.score(X_train_pca,y_train_pca)
    test_accuracy[i] = knn.score(X_test_pca, y_test_pca)

plt.plot(neighbors, test_accuracy, label='Testing Accuracy')
plt.plot(neighbors, train_accuracy, label='Training Accuracy')
plt.legend()
plt.show()
```



High accuracy when k=1

```
In [49]: knn = KNeighborsClassifier(n_neighbors=1)
knn.fit(X=X_train_pca, y=y_train_pca)
knn.score(X_test_pca, y_test_pca)
```

Out[49]: 0.8333333333333334

```
In [50]: from sklearn.metrics import confusion_matrix
y_pred_pca = knn.predict(X_test_pca)
confusion_matrix(y_test_pca, y_pred_pca)
```

Out[50]: array([[18, 2],
[5, 17]], dtype=int64)

```
In [51]: from sklearn.metrics import classification_report
print(classification_report(y_test_pca, y_pred_pca))
```

	precision	recall	f1-score	support
M	0.78	0.90	0.84	20
R	0.89	0.77	0.83	22
accuracy			0.83	42
macro avg	0.84	0.84	0.83	42
weighted avg	0.84	0.83	0.83	42

Accuracy: 83.3%

7 out of 42 wrong predictions

Decision Tree

```
In [52]: from sklearn.tree import DecisionTreeClassifier
         clf = DecisionTreeClassifier()
         clf.fit(X_train_pca, y_train_pca)
         clf.score(X_test_pca, y_test_pca)
```

Out[52]: 0.7619047619047619

```
In [53]: from sklearn.metrics import confusion_matrix
         y_pred_pca = clf.predict(X_test_pca)
         confusion_matrix(y_test_pca, y_pred_pca)
```

Out[53]: array([[15, 5],
 [5, 17]], dtype=int64)

```
In [54]: from sklearn.metrics import classification_report
         print(classification_report(y_test_pca, y_pred_pca))
```

	precision	recall	f1-score	support
M	0.75	0.75	0.75	20
R	0.77	0.77	0.77	22
accuracy			0.76	42
macro avg	0.76	0.76	0.76	42
weighted avg	0.76	0.76	0.76	42

Accuracy: 76.19%

10 out of 42 wrong predictions

SVM

```
In [55]: from sklearn.svm import SVC
         svm = SVC()
         svm.fit(X_train_pca, y_train_pca)
         svm.score(X_test_pca, y_test_pca)
```

C:\Users\SR1407SM1106\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
"avoid this warning.", FutureWarning)

Out[55]: 0.7857142857142857

```
In [56]: from sklearn.metrics import confusion_matrix
         y_pred_pca = svm.predict(X_test_pca)
         confusion_matrix(y_test_pca, y_pred_pca)
```

Out[56]: array([[18, 2],
 [7, 15]], dtype=int64)

```
In [57]: from sklearn.metrics import classification_report  
print(classification_report(y_test_pca, y_pred_pca))
```

	precision	recall	f1-score	support
M	0.72	0.90	0.80	20
R	0.88	0.68	0.77	22
accuracy			0.79	42
macro avg	0.80	0.79	0.78	42
weighted avg	0.81	0.79	0.78	42

Accuracy: 78.57%

9 out of 42 wrong predictions

Algorithm	Accuracy (Normal)	Accuracy (PCA)	Wrong Predictions (Normal)	Wrong Predictions (PCA)
KNN	78.5	83.3	9	7
DT	90.47	76.19	4	10
SVM	64.28	78.57	15	9

PCA worked well for KNN and SVM. Decision Tree gave highest accuracy


Clustering:

K means clustering

```
In [90]: from sklearn.cluster import KMeans
wcss = []

for k in range(1, 15):
    kmeansForLoop = KMeans(n_clusters = k)
    kmeansForLoop.fit(X)
    wcss.append(kmeansForLoop.inertia_)

plt.figure(figsize = (10, 10))
plt.plot(range(1, 15), wcss)
plt.xlabel("K value")
plt.ylabel("WCSS")
plt.show()
```



```
In [76]: from sklearn.cluster import KMeans
df = data
df.head()
```

Out[76]:

	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	0.2111	...	0.0027	0.0065	0.0159
0	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	0.2872	...	0.0084	0.0089	0.0048
1	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	0.6194	...	0.0232	0.0166	0.0095
2	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	0.1264	...	0.0121	0.0036	0.0150
3	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	0.4459	...	0.0031	0.0054	0.0105
4	0.0286	0.0453	0.0277	0.0174	0.0384	0.0990	0.1201	0.1833	0.2105	0.3039	...	0.0045	0.0014	0.0038

5 rows × 61 columns

```
In [77]: kmeans = KMeans(n_clusters=2)
```

```
In [78]: kmeans.fit(df.drop('R',axis=1))
```

```
Out[78]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
               n_clusters=2, n_init=10, n_jobs=None, precompute_distances='auto',
               random_state=None, tol=0.0001, verbose=0)
```

```
In [80]: kmeans.cluster_centers_
```

```
Out[80]: array([[0.03000612, 0.03953673, 0.04685204, 0.06019796, 0.08832143,
                  0.11321122, 0.13089592, 0.14420918, 0.18996531, 0.23429592,
                  0.2737102 , 0.28927551, 0.33128163, 0.38904796, 0.44173878,
                  0.5304      , 0.60288469, 0.65269592, 0.70914184, 0.74849388,
                  0.78412755, 0.79563673, 0.77186939, 0.74666939, 0.7076602 ,
                  0.69397857, 0.67020816, 0.60677653, 0.51466837, 0.44462653,
                  0.37925102, 0.32470408, 0.30264796, 0.28699694, 0.28031735,
                  0.27957857, 0.28383265, 0.26861122, 0.26827755, 0.26862653,
                  0.23407653, 0.22373673, 0.20443673, 0.18207449, 0.15821122,
                  0.12247857, 0.1053      , 0.07932959, 0.04807551, 0.0196898 ,
                  0.01580306, 0.01308776, 0.01093571, 0.01158673, 0.0095898 ,
                  0.00877041, 0.00800408, 0.0085      , 0.00769082, 0.00622041],
                [0.02849083, 0.03745963, 0.04112661, 0.04852752, 0.06322202,
                  0.09685596, 0.11322569, 0.12610642, 0.16602844, 0.18482385,
                  0.20280917, 0.21595229, 0.22163394, 0.21555046, 0.21326147,
                  0.24329083, 0.2489156 , 0.27356055, 0.32107431, 0.3970789 ,
                  0.4519422 , 0.47128165, 0.53664954, 0.60718807, 0.64647982,
                  0.70569541, 0.73080183, 0.7714211 , 0.75628257, 0.70526514,
                  0.62049174, 0.54347706, 0.51935872, 0.50451468, 0.48926606,
                  0.47521927, 0.43290826, 0.40105596, 0.37597064, 0.34982752,
                  0.34104495, 0.32729725, 0.28406789, 0.2409055 , 0.23170183,
                  0.19513578, 0.13803394, 0.10190459, 0.05551835, 0.02097431,
                  0.0162422 , 0.01381743, 0.01054404, 0.01031468, 0.00904037,
                  0.00765046, 0.00756147, 0.00744954, 0.00815688, 0.00679541]])
```

```
In [84]: def converter(cluster):
           if cluster=='R':
               return 1
           else:
               return 0
```

```
In [85]: df['Cluster'] = df['R'].apply(converter)
df.head()
```

```
Out[85]:
```

	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	0.2111	...	0.0065	0.0159	0.007%
0	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	0.2872	...	0.0089	0.0048	0.0094
1	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	0.6194	...	0.0166	0.0095	0.0180
2	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	0.1264	...	0.0036	0.0150	0.0085
3	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	0.4459	...	0.0054	0.0105	0.0110
4	0.0286	0.0453	0.0277	0.0174	0.0384	0.0990	0.1201	0.1833	0.2105	0.3039	...	0.0014	0.0038	0.0015

5 rows × 62 columns

```
In [88]: kmeans.labels_
```

```
Out[88]: array([0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1,
        1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1,
        1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0,
        1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0,
        0, 0, 0, 1, 1, 1, 1, 1, 1])
```

```
In [89]: from sklearn.metrics import confusion_matrix, classification_report
print(confusion_matrix(df['Cluster'], kmeans.labels_))
print(classification_report(df['Cluster'], kmeans.labels_))
```


```
[[58 53]
 [40 56]]
```

	precision	recall	f1-score	support
0	0.59	0.52	0.56	111
1	0.51	0.58	0.55	96
accuracy			0.55	207
macro avg	0.55	0.55	0.55	207
weighted avg	0.56	0.55	0.55	207


Accuracy: 55%

93 wrong predictions out of 207

```
In [92]: plt.figure(figsize = (10, 7))
plt.scatter(df['0.0200'][df["Cluster"] == 0], df["0.0371"][df["Cluster"] == 0],
            color = "red")
plt.scatter(df["0.0200"][df["Cluster"] == 1], df["0.0371"][df["Cluster"] == 1],
            color = "green")
plt.xlabel('0.0200')
plt.ylabel('0.0371')
plt.show()
```




```
In [95]: # plotting centroids
clusters = kmeans.fit_predict(X)
plt.figure(figsize = (10, 7))
plt.scatter(df["0.0200"], df["0.0371"], c = clusters, alpha = 0.5)
plt.scatter(kmeans.cluster_centers[:, 0], kmeans.cluster_centers[:, 1], color = "red", alpha = 1)
plt.xlabel('0.0200')
plt.ylabel('0.0371')
plt.show()
```



Hierarchical Clustering

```
In [97]: df_X = df.drop('R', axis=1)
from scipy.cluster.hierarchy import linkage, dendrogram
merg = linkage(df_X, method='ward')
dendrogram(merg, leaf_rotation = 90)
plt.xlabel('Data points')
plt.ylabel('Euclidean distance')
plt.show()
```




```
In [98]: from sklearn.cluster import AgglomerativeClustering
hc = AgglomerativeClustering(n_clusters = 2, affinity = "euclidean", linkage = "
ward")
cluster = hc.fit_predict(df_X)
df_X["label"] = cluster
```

```
In [102]: df_X.label.value_counts()
```

```
Out[102]: 0    142
          1     65
          Name: label, dtype: int64
```

```
In [104]: plt.figure(figsize = (10, 7))
plt.scatter(df_X['0.0200'][df_X["label"] == 0], df_X["0.0371"][df_X["label"] ==
0], color = "red")
plt.scatter(df_X["0.0200"][df_X["label"] == 1], df_X["0.0371"][df_X["label"] ==
1], color = "green")
plt.xlabel('0.0200')
plt.ylabel('0.0371')
plt.show()
```




```
In [107]: def converter(cluster):
            if cluster=='R':
                return 1
            else:
                return 0

df_X['Cluster'] = df['R'].apply(converter)

from sklearn.metrics import confusion_matrix, classification_report
print(confusion_matrix(df_X['Cluster'], hc.labels_))
print(classification_report(df_X['Cluster'], hc.labels_))
```

```
[[72 39]
 [70 26]]
```

	precision	recall	f1-score	support
0	0.51	0.65	0.57	111
1	0.40	0.27	0.32	96
accuracy			0.47	207
macro avg	0.45	0.46	0.45	207
weighted avg	0.46	0.47	0.45	207

Accuracy: 47%

109 wrong predictions out of 207

DBSCAN

```
In [113]: from sklearn.cluster import DBSCAN
            from sklearn.preprocessing import StandardScaler
X.head()
```

```
Out[113]:
```

	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	0.2111	...	0.0232	0.0027	0.0061
0	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	0.2872	...	0.0125	0.0084	0.0081
1	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	0.6194	...	0.0033	0.0232	0.0161
2	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	0.1264	...	0.0241	0.0121	0.0031
3	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	0.4459	...	0.0156	0.0031	0.0051
4	0.0286	0.0453	0.0277	0.0174	0.0384	0.0990	0.1201	0.1833	0.2105	0.3039	...	0.0104	0.0045	0.0011

5 rows × 60 columns

```
In [109]: X_ss = StandardScaler().fit_transform(X)
```

```
In [110]: db = DBSCAN().fit(X_ss)
            cluster_labels = db.labels_
            num_clusters = len(set(cluster_labels))
```

```
Out[112]: 1
```

In []: